

Dr. Alexander Wright

Staff AI Research & LLM Systems Engineer

alex.wright@stanford-alumni.org · (650) 555-0101 · Palo Alto, CA · github.com/alexwright-ai

PROFESSIONAL SUMMARY

Staff AI Engineer with 8 years experience architecting large language model training and low-latency inference runtimes. Author of 4 top-tier conference publications (NeurIPS, ICML). Creator of an open-source distributed LLM serving engine with 14k+ GitHub stars. Expert in PyTorch, CUDA, vLLM, TensorRT-LLM, and vector search infrastructure.

CORE TECHNICAL COMPETENCIES

PyTorch ? LLM ? CUDA ? vLLM ? Transformers ? Python ? C++ ? TensorRT-LLM ? FAISS ? Vector Search ? RAG ? Distributed Systems

WORK EXPERIENCE & IMPACT

Staff AI Systems Engineer · Cognitive Scale AI Labs (2021 - Present)

- Architected high-throughput distributed inference cluster serving 70B parameter models at 18ms TTFT (Time To First Token) for 2M daily requests.
- Engineered custom CUDA kernel kernels and PagedAttention layers, reducing GPU memory footprint by 44% and saving \$1.8M annually in cloud costs.
- Led team of 6 research engineers delivering multimodal retrieval-augmented generation (RAG) platform with hybrid dense-sparse vector indexing.

Senior Machine Learning Engineer · Google Brain / Deep Research (2018 - 2021)

- Developed distributed transformer training pipelines scaling across 512 TPU v3 cores with model parallelism.
- Published 2 first-author papers at NeurIPS on sample-efficient sequence distillation.

EDUCATION & CREDENTIALS

Ph.D. in Computer Science (Artificial Intelligence) · Stanford University (2018) · B.S. in Math & CS (Highest Honors) · MIT